

# Geomatic Techniques to Support Phytosanitary Products Tests within the EPPO Standard Framework

Samuele Bumbaca

University of Turin

August 28, 2025



**SAGEA**  
Group



**DISAFA**



**UNIVERSITÀ  
DI TORINO**

# The Traditional Approach to Agricultural Trials

Block 3	R	C	T
Block 2	T	R	C
Block 1	C	T	R

C Control  
T Tested Product  
R Reference Product

## ANOVA Model:

$$y_{ij} = \mu + \alpha_i + \beta_j + \varepsilon_{ij}$$

Where:

- $y_{ij}$  = response
- $\mu$  = overall mean
- $\alpha_i$  = treatment effect
- $\beta_j$  = block effect
- $\varepsilon_{ij}$  = random error

### Note:

This is the **additive model**. Modern approaches may include interaction

terms:  $\alpha_i \times \beta_j$

# Key Assumptions of Traditional ANOVA


## Statistical Assumptions:

- **Randomization:** Treatments randomly assigned within blocks
- **Replication:** Each treatment appears in each block
- **Independence:** Observations are independent given the design
- **Homoscedasticity :** Equal variances across treatments
- **Normality:** Residuals follow normal distribution


## Consequences of Assumption Violations:

- **Invalid conclusions of parametric tests:** Need for non-parametric tests leading to reduced statistical power

# The Right Blocking: Capturing Environmental Variability



Environmental Gradient:




Variability

C Control  
T Tested Product  
R Reference Product


## Success of Blocking Strategy:

- **Within-block homogeneity:** Treatments compared under similar conditions
- **Between-block heterogeneity:** Environmental gradient captured by block effects

# The Wrong Blocking: Assumption Violation



Environmental Gradient:



Variability

C Control  
T Tested Product  
R Reference Product

## Heteroscedasticity Assumption Violation Problem:

- **Blocks fail to capture environmental variability:** Treatments compared under different conditions
- **Invalid parametric test:** Residual variance differs across treatments

# Current Limitations in Statistics for Agricultural Trials

---

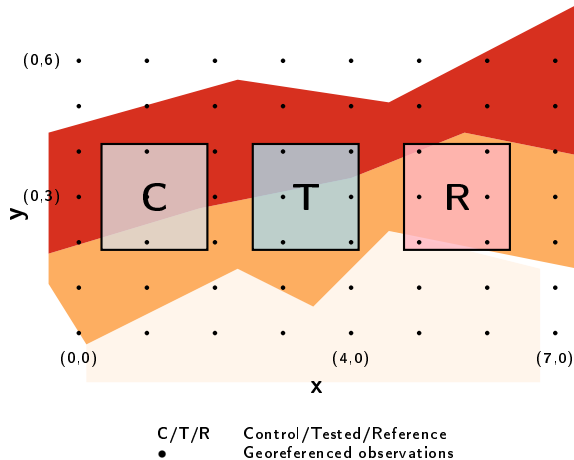
## Traditional Approach Issues:

- **Human-dependent blocking:** Environmental variability assessment relies on experimenter experience
- **A priori identification:** Must identify variance sources BEFORE data collection

## The Challenge:

*How can we capture environmental variability mathematically rather than through human judgment?*

# Geostatistical Approach: Spatial Linear Mixed Models



## Spatial LMM:

$$y(s_i) = \mu + \alpha_j + f(s_i) + \varepsilon_i$$

Where:

- $y(s_i)$  = response at  $s_i$
- $\mu$  = overall mean
- $\alpha_j$  = treatment effect
- $f(s_i)$  = spatial random field
- $\varepsilon_i$  = error
- $s_i = (x_i, y_i)$  = coordinates

## Benefits:

- **No blocking:** Spatial correlation captures variability
- **Post-hoc:** No a priori variance identification
- **Homoscedasticity:** Assumption satisfied in more cases in respect blocking

# Statistical Methods Comparison: Introduction

## Comparison Objective:

Evaluate the performance of **traditional RCBD** versus **spatial geostatistical methods** (SpATS) in capturing environmental variability and estimating treatment effects.

## Synthetic Dataset:

- **54 observations**(6×9 grid)
- **3 treatments**: Control (0 t/ha), Reference (0.5 t/ha), Test (1.0 t/ha)
- **3 blocks**(18 plots each)
- **Environmental zones**: Low (-1.5 t/ha), Medium (0 t/ha), High (+1.5 t/ha)

## Tested Models:

- 1 **RCBD Model**: Linear Mixed Model with random block effects

$$y_{ij} = \mu + \alpha_i + \beta_j + \varepsilon_{ij}$$

- 2 **SpATS Model**: Spatial model with PSANOVA splines

$$y(s) = \mu + \alpha_i + f(s) + \varepsilon(s)$$

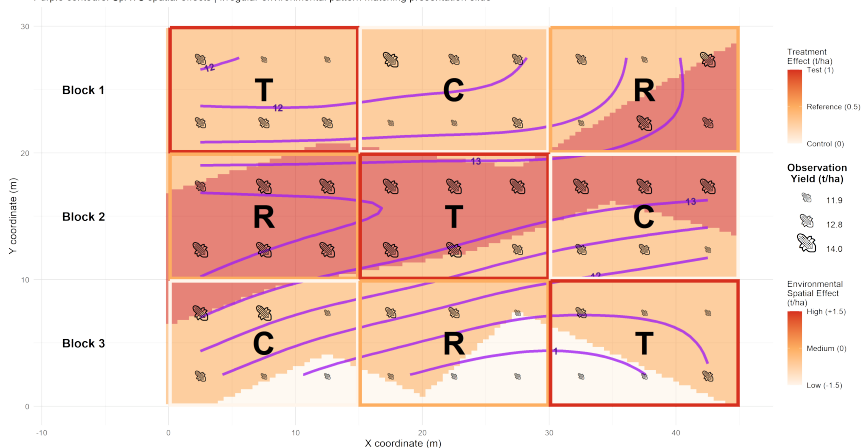
Where:  $\alpha_i$  = treatment effects,  $\beta_j$  = block effects,  $f(s)$  = spatial smooth



# Statistical Methods Comparison: The Field Trial Design

## Irregular Environmental Gradient Trial Design

Purple contours: SpATS spatial effects | Irregular environmental pattern matching presentation slide



# Statistical Methods Comparison: Results

## Model Performance (Mean Absolute Errors tonn/ha):

Model	Treat. Error	Env. Error
RCBD Model	0.13	0.62
SpATS Spatial	0.03	0.45

## Treatment Effect Estimation (tonn/ha):

Treatment	True	RCBD	SpATS
Control	0.00	0.00	0.00
Reference	0.50	0.40	0.45
Test	1.0	0.69	0.94

## Key Findings:

- Both models satisfied assumptions
- SpATS outperformed RCBD:
  - 3.8× better treatment effect estimation
  - 1.4× better environmental effect estimation
- RCBD underestimated by 20-31%
- SpATS <6% error

## Implications:

Even when traditional RCBD meets statistical assumptions, **spatial modeling provides superior accuracy** in treatment effect estimation by properly accounting for environmental spatial variability.

# The Missing Link: Spatial Coordinates

---

## Geostatistical Methods Advantages:

- ✓ **Mathematical modeling** of environmental variability
- ✓ **Post-hoc analysis** - no need for prior knowledge of the environment variables and of their distribution
- ✓ **Superior performance** in handling spatial heterogeneity
- ✓ **EPPO recognized** approach (PP1/152(4) - Design and analysis of efficacy evaluation trials)

## Current Barrier:

- ✗ **Requires spatially referenced observations**
- ✗ **Traditional manual assessments lack coordinates**
- ✗ **Implementation gap** in practical field trials

# Central Research Question

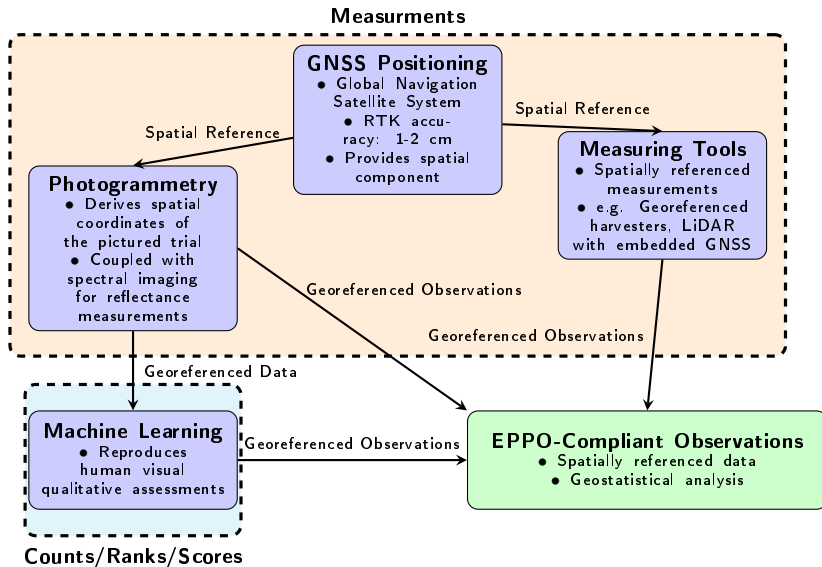
---

**Can geomatics technologies provide spatially referenced observations that enable geostatistical analysis within EPPO-compliant Plant Protection Product trials?**

## Specific Objectives:

- 1 Establish which geomatics technologies can be used to collect spatially referenced observations
- 2 Demonstrate the feasibility of collect spatially referenced observations in compliant with EPPO standards
- 3 Validate performance against traditional methods
- 4 Provide practical implementation guidelines

# Geomatic Technologies: Workflow for Spatially Referenced Observations



# Georeferencing EPPO Standard Assessments

Table: EPPO's types of variables

Type of Variable	Measurement	Ranking	Scoring
Continuous not limited	X		
Continuous limited	X		
Discrete	X		
Ordinal		X	X
Nominal			X
Binary			X

Summary from EPPO PP 1/152: Design and analysis of efficacy evaluation trials

## Current State of Georeferencing in Agricultural Trials:

Tool-based measurements (e.g., yield harvesters) can be easily georeferenced by integrating GNSS receivers on the tool. For visual assessments as counting, scoring or ranking, a method to transform georeferenced data to georeferenced observations is needed.

# Machine Learning Inference on Georeferenced Data

---

## **Machine Learning for Visual Assessments:**

Machine Learning (ML) can reproduce human visual assessments, allowing for spatially referenced observations.



## **Machine Learning Limitations:**

ML needs training data to reproduce human assessments. Understand data requirement to fit EPPO standard is critical.



## **Application Case Studies:**

This thesis provides a case study to prove the possibility to get georeferenced observations for each variable type that need for a ML step.

## EPPO PP 1/333(1): Digital Technologies in PPP Trials

ML integrated assessments must meet the same quality standards as manual assessments and require validation through comparison with manual assessments (golden sample).

### Validation Benchmarks<sup>a</sup>

<sup>a</sup>Based on EPPO PP 1/333(1): Use of digital technologies in efficacy and selectivity trials

- **Continuous/Discrete:**  $R^2 > 0.85$  (1:1 relationship)
- **Ordinal/Nominal:** Cohen's  $\kappa > 0.7$
- **Binary:** Accuracy  $> 0.85$



# Georeferencing Gap in EPPO Standard Assessments

	Type of Variable	Measurement	Ranking	Scoring
	Continuous not limited	X		
	Continuous limited	X		
→	Discrete	X		
→	Ordinal		X	X
→	Nominal			X
→	Binary			X

## Case Studies:

This thesis aim to prove the reliability of georeferencing every EPPO standard assessment Each case study addresses a specific variable type as defined in the EPPO standards

- **Discrete (Counts)** : Plant counting
- **Ordinal** : Phytotoxicity scoring
- **Nominal** and **Binary** : Disease detection

# Georeferencing Counts (Discrete Variable)

	Type of Variable	Measurement	Ranking	Scoring
	Continuous not limited	X		
	Continuous limited	X		
→	Discrete	X		
	Ordinal		X	X
	Nominal			X
	Binary			X

## Georeferencing Counts:

- **Counts** are discrete variables required for measuring density of individuals (e.g. plant density in PP1/46 (3) - Wireworms).
- the **Case Study**: Counting plants from georeferenced photogrammetric orthomosaics by ML Object Detection.
- this study is discussed in the scientific article **Bumbaca, S.; Borgogno-Mondino, E.C. On the Minimum Dataset Requirements for Fine-Tuning an Object Detector for Arable Crop Plant Counting: A Case Study on Maize Seedlings. Remote Sens. 2025, 17, 2190. DOI: 10.3390/rs1713219061**

# Plant Counting: The Core Challenge

## The Critical Need after EPPO Assessments:

- Plant counting is **fundamental** in precision agriculture and plant breeding
- Traditional manual counting is **time-consuming** and **variable**
- Computer vision offers automation and standardization but requires **dataset size and quality** characterization to prove the reliability for this task.

## EPPO Benchmark Standards:

Coefficient of determination ( $R^2$ )  $\geq 0.85$  w.r.t. manual counting

## Research Gap:

*What are the minimum dataset requirements to achieve these benchmarks across different inference datasets?*

# Case Study: Maize Seedling Counting

## Why Maize Seedlings?

- **Data availability:** Most represented plant in scientific <sup>1a 2b</sup> and public datasets <sup>3c 4d</sup>
- **Economic importance:** Most important crop worldwide by production <sup>5e</sup>
- **Optimal detection conditions:** Low overlapping and fixed spacing at V3-V5 stage

---

<sup>a</sup>David et al. Plant detection and counting from high-resolution RGB images acquired from UAVs. *bioRxiv*, 2021

<sup>b</sup>Liu et al. IntegrateNet: A deep learning network for maize stand counting from UAV imagery. *IEEE Geosci. Remote Sens. Lett.*, 19:6512605, 2022

<sup>c</sup>Maize\_seeding dataset.

[https://universe.roboflow.com/objectdetection-hytat/maize\\_seeding](https://universe.roboflow.com/objectdetection-hytat/maize_seeding)

<sup>d</sup>Maize-seedling-detection dataset.

<https://universe.roboflow.com/fyxdds-icloud-com/maize-seedling-detection>

<sup>e</sup>FAO. *Agricultural Production Statistics 2010–2023*. FAOSTAT, Rome, Italy, 2024

## Growth Stage Selection:

V3-V5 (BBCH 13-15) <sup>6a</sup>

Broader Applications:

## Handcrafted Methods (HC):

- **Traditional approach:** Still used in agricultural applications <sup>1a 2b</sup>
- **Color thresholding** in HSV space + morphological operations
- **Trade-off:** High precision on subset vs. limited generalizability

<sup>a</sup>David et al. Plant detection and counting from high-resolution RGB images acquired from UAVs. *bioRxiv*, 2021

<sup>b</sup>García-Martínez et al. Digital count of corn plants using UAVs and cross correlation. *Agronomy*, 10:469, 2020

## CNN-based Models:

### Convolutional Neural Networks <sup>3a</sup>

- **YOLO family:** YOLOv5, YOLOv8
- **Faster R-CNN** <sup>4b</sup>
- Grid-like image processing

## Transformer-based:

### Attention mechanisms <sup>5a</sup>

- **DETR** <sup>6b</sup>
- **RT-DETR:** Real-time performance
- Sequence of patches processing
- Superior with scarce data <sup>7c</sup>

# Representative Architectures: YOLO vs RT-DETR

## YOLO Family (Pure CNN):

### Why chosen as representatives:

- **Large adoption** in agriculture <sup>1a</sup>
- **Good precision** and low dataset requirements vs. other CNNs <sup>2b 3c</sup>
- **Speed-accuracy** optimizations

### Variants tested:

- **YOLOv5**: CSPDarknet53 backbone
- **YOLOv8**: Improved architecture with decoupled head

---

<sup>a</sup>Badgujar et al. Agricultural object detection with YOLO algorithm. *Comput. Electron. Agric.*, 223:109090, 2024

<sup>b</sup>Tan et al. EfficientDet: Scalable efficient object detection, 2020

## RT-DETR (Transformer-mixed):

### Why chosen:

- **Outperforms** YOLOv5 and YOLOv8 <sup>4a</sup>
- **Real-time** transformer detector
- **Hybrid approach**: CNN backbone + transformer decoder

### Architecture advantages:

- Enhanced feature extraction
- Better handling of spatial relationships
- Improved performance with limited data

---

<sup>a</sup>Zhao et al. DETRs beat YOLOs on agricultural object detection. *Yi*, 2304.00060

# The Dataset Requirements Challenge in Agriculture

## Critical Research Gap:

Few studies focus on minimum dataset requirements for robust plant detectors <sup>1a</sup>  
<sup>2b</sup>, despite extensive work on agricultural object detection <sup>3c 4d 5e</sup>

<sup>a</sup>David et al. Plant detection and counting from high-resolution RGB images acquired from UAVs. *bioRxiv*, 2021

<sup>b</sup>Andvaag et al. Counting canola: Toward generalizable aerial plant detection. *Plant Phenomics*, 6:0268, 2024

<sup>c</sup>Barreto et al. Automatic UAV-based counting of seedlings. *Comput. Electron. Agric.*, 191:106493, 2021

<sup>d</sup>Jiang et al. Deep seedling: Deep convolutional network, 2019

<sup>e</sup>Liu et al. IntegrateNet: A deep learning network for maize stand counting from UAV imagery. *IEEE Geosci. Remote Sens. Lett.*, 19:6512605, 2022

## Known Performance Factors:

- **Dataset size:** Performance directly related to training data amount <sup>6a</sup>
- **Data quality:** Annotation accuracy critically affects model performance <sup>7b</sup>
- **Model architecture:** Different models require different dataset sizes for same performance <sup>8c 9d</sup>

# Object Detection Paradigms for Plant Counting

---

## Many-Shot Models:

- **CNN-based:** YOLOv5, YOLOv8
- **Transformer-mixed:** RT-DETR, YOLO11
- Require extensive labeled datasets
- State-of-the-art performance

## Zero-Shot Models:

- **OWLv2:** Open-vocabulary detection
- Vision-language foundation models
- Text prompt-based detection
- No training data required

## Few-Shot Models:

- **CD-ViTO:** Cross-domain adaptation
- Meta-learning approaches
- 1-50 training examples
- Promising but unvalidated for agriculture

## Handcrafted Methods:

- **Color thresholding** + agronomic knowledge
- High precision in constrained scenarios
- Still used as baseline/annotation tool
- Limited generalizability



# Study Objectives and Research Questions

## Primary Objective:

Determine minimum dataset size and quality required to achieve EPPO benchmarks ( $R^2 \geq 0.85$ ) for maize seedling detection across different object detection paradigms.

## Specific Research Questions:

- 1 What is the impact of **dataset source** (in-domain vs. out-of-distribution)?
- 2 How do **model architectures** affect dataset requirements?
- 3 What is the minimum acceptable **annotation quality**?
- 4 Can **few-shot/zero-shot** approaches meet agricultural benchmarks?
- 5 What role do **handcrafted methods** play in the DL era?

## Case Study Focus:

**Maize seedlings** (*Zea mays* L.) at V3-V5 growth stage from georeferenced orthomosaics

# Dataset Collection and Preparation

## Dataset Sources:

### Out-of-Distribution (OOD):

- Scientific literature: 778 tiles
- Internet repositories: 390 tiles
- Pre-annotated datasets

### In-Domain (ID):

- 3 study sites: 450 training + 60 test tiles
- Phantom 4 Pro v2.0 @ 10m AGL
- Bundle adjustment error: 38mm (GNSS VRS-NRTK)

## Technical Specs:

- **Resolution:** 5 mm/pixel
- **Tile size:** 224×224 pixels
- **Coverage:** 1.12×1.12 meters
- **Content:** 2 maize rows per tile
- **Annotation:** Squared bounding boxes centered on stems

## Key Insight:

Tile size optimized for row pattern identification and model compatibility

# Handcrafted Object Detector: Two-Stage Pipeline

## Stage 1 - HC1 (Detection):

- 1 **Color thresholding** in HSV space
- 2 **Connected components** analysis
- 3 **Size filtering** based on leaf area
- 4 Outputs: Potential plant regions

*Result: High recall, many false positives*

## Stage 2 - HC2 (Verification):

- 1 **RANSAC** line fitting for row detection
- 2 **Row spacing** validation
- 3 **Plant count** verification per row
- 4 **Agronomic knowledge** application

*Result: High precision, limited coverage*

## Algorithm Performance:

Dataset	R <sup>2</sup>	Coverage
ID_1	0.95	7.8%
ID_2	0.93	4.2%
ID_3	0.87	1.8%

## Trade-off:

Excellent accuracy on subset of data  
vs. limited generalizability

# Deep Learning Model Configuration

## Many-Shot Models (Ultralytics Implementation):

### CNN-based:

- YOLOv5 (n, s, m, l, x)
- YOLOv8 (n, s, m, l, x)

### Transformer-mixed:

- YOLO11 (n, s, m, l, x)
- RT-DETR (l, x)

### Training Settings:

- Batch size: 16
- Max epochs: 200
- Early stopping: 15 epochs
- Mixed precision training
- Default Ultralytics augmentation

## Few-Shot: CD-ViT0

- ViT-S/B/L backbones (22M/86M/307M params)
- 1, 5, 10, 30, 50 shots tested
- Cross-domain adaptation

## Zero-Shot: OWLv2

- ViT-B/16, ViT-L/14 encoders
- Base, fine-tuned, ensemble variants
- 11 different text prompts tested

# Experimental Design and Evaluation Metrics

## Dataset Size Investigation:

- **Many-shot:** 10-150 images in steps of 10
- **Few-shot:** 1, 5, 10, 30, 50 shots
- **Zero-shot:** No training data required

## Dataset Quality Investigation:

Annotation quality reduced from 100% to 10% in 10% steps for successful models

## Evaluation Metrics:

### Counting Performance:

- $R^2$  (coefficient of determination)
- $RMSE$  (root mean square error)
- $MAPE$  (mean absolute percentage error)

### Detection Performance:

### Performance Modeling:

Empirical functions tested:

- $f(x) = a \ln(x) + b$
- $f(x) = a \arctan(bx) + c$
- $f(x) = ax^{1/b} + c$

Best fit selected by  $R^2$  (GoF)

# Testing Protocol and Infrastructure

## Hardware Configuration:

- **CPU:** Intel Xeon E5-2670 v3 @ 2.30GHz
- **RAM:** 64.0 GB
- **GPU:** NVIDIA RTX A5000 (24GB VRAM)
- **Implementation:** Ultralytics, HuggingFace Transformers

## SAHI Testing Method:

- ① **Slice:** Test images into overlapping patches
- ② **Detect:** Run model on each patch
- ③ **Merge:** Combine predictions with NMS
- ④ **Threshold:** Apply confidence score filtering

*Rationale: Handles object occlusion at tile boundaries*

## Confidence Thresholds:

0, 0.05, 0.1, 0.15, 0.2, 0.25, 0.29, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 0.95, 0.99

**Best  $R^2$  selected** across all thresholds for each model

## Key Principle:

# Critical Impact of Dataset Source

## Major Finding:

**NO** out-of-distribution model achieved benchmark performance ( $R^2 \geq 0.85$ )

## OOD Results:

- **Best  $R^2$ :**  $< 0.5$  (all models)
- **Best MAPE:** 20
- **GoF values:**  $< 0.2$  (poor predictability)
- **Dataset size:** Up to 1,168 images tested

## Domain Gap Challenge:

- Environmental conditions
- Lighting variations

## ID Success Stories:

### Models achieving $R^2 \geq 0.85$ :

- YOLOv5n: 130 samples
- YOLOv5s: 130 samples
- YOLOv8n: 110 samples
- RT-DETR L: 60 samples
- RT-DETR X: 100 samples

**GoF values:**  $> 0.3$  (high predictability)

# Architecture-Specific Dataset Requirements

## CNN-based Models (YOLO family):

- **YOLOv5n**: 130 samples (1.9M params)
- **YOLOv5s**: 130 samples (7.2M params)
- **YOLOv8n**: 110 samples (3.2M params)

*Pattern: Larger models → more samples needed*

## Transformer-mixed Models:

- **RT-DETR L**: 60 samples (60M params)
- **RT-DETR X**: 100 samples (80M params)

*RT-DETR L most efficient*

## Key Insights:

- **Transformers more sample-efficient** than CNNs
- **Long-range dependencies** better captured
- **Trade-off**: Higher computational cost

## Practical Decision:

**CNN approach**: Collect 130 samples + lower compute

**Transformer approach**: Collect 60 samples + higher compute

## Performance Predictability:

Logarithmic relationship between dataset size and performance enables resource planning ( $GoF > 0.3$  for successful models)



# Dataset Quality Requirements

## Quality Tolerance Analysis:

Models achieving benchmark with reduced annotation quality:

## Successful Quality Reductions:

- **YOLOv5n**: 85% quality (130 samples)
- **YOLOv5s**: 90% quality (130 samples)
- **YOLOv8n**: 85% quality (110 samples)
- **RT-DETR X**: 65% quality (100 samples)

## RT-DETR L Sensitivity:

Failed to maintain benchmark with ANY quality reduction (60 samples baseline)

## Quality-Quantity Relationship:

**Key Finding:** Smaller datasets are more sensitive to annotation errors

**RT-DETR L:** Minimal dataset (60 samples) → Each annotation critical

**RT-DETR X:** Larger dataset (100 samples) → Error redundancy tolerance

## Practical Strategy:

**Option 1:** Perfect annotations + minimal dataset

**Option 2:** Good quality annotations + larger dataset

**Option 3:** Semi-automated annotation workflows

# Few-Shot and Zero-Shot: Current Limitations

## Major Finding:

Neither few-shot nor zero-shot approaches achieved benchmark performance

## Few-Shot Results (CD-ViTO):

### Best Performance (ViT-B, 50 shots):

- **RMSE:** 3.9 (vs. benchmark 0.39)
- **MAPE:** 25
- **mAP:** 0.5
- **Error rate:** 1 plant in 4 miscounted

*Pattern: Performance plateaus after 30 shots*

## Zero-Shot Results (OWLv2):

- **$R^2$ :** Always  $< 0$  (worse than mean prediction)
- **RMSE:** 5-25 (extremely high)
- **MAPE:** 40-140%
- **Prompt sensitivity:** High variability across 11 prompts

## Why the Poor Performance?

# The Value of Hybrid Approaches

## Handcrafted Method Performance:

### Strengths:

- $R^2 = 0.87-0.95$  (excellent accuracy)
- $RMSE = 0.11-0.18$  (below benchmark)
- Domain knowledge integration
- High precision when applicable

### Limitations:

- Coverage: 1.8-7.8% of tiles only
- Color-thresholding bias
- Limited generalizability

## Hybrid Strategy Potential:

### 1. Bootstrap Training:

- HC method generates high-quality annotations
- Deep learning models trained on HC output
- Overcomes manual annotation bottleneck

### 2. Quality Filtering:

- OOD/few-shot/zero-shot occasional good predictions
- HC2 validates agronomic patterns
- Reduces color-thresholding bias

## Future Work Direction:

# Practical Implementation Guidelines

## Resource Optimization Strategy:

**Step 1:** Focus on minimum viable dataset size (60-130 images) **Step 2:** Logarithmic relationship → diminishing returns beyond minimum **Step 3:** Quality vs. quantity trade-off consideration

## Implementation Pathways:

### High-Resource Scenario:

- RT-DETR L + 60 perfect annotations
- Higher computational investment
- Fastest deployment

### Medium-Resource Scenario:

- YOLOv8n + 110 good annotations
- Balanced compute/annotation effort
- Robust performance

### Low-Resource Scenario:

- YOLOv5n + 130 annotations (85% quality)
- Semi-automated annotation

## Critical Success Factors

- ❶ **In-domain data:** Non-negotiable requirement
- ❷ **Architecture choice:** Based on resource constraints
- ❸ **Quality assessment:** Monitor annotation accuracy
- ❹ **Validation protocol:** SAHl testing recommended

## Industry Adoption

## Core Findings

- **In-domain training data is mandatory** - OOD approaches fail to achieve benchmarks
- **Architecture matters**: Transformer-mixed models (RT-DETR) require 50% fewer samples than CNN-based models
- **Quality tolerance exists**: Models maintain performance with 65-90% annotation quality
- **Current limitations**: Few-shot and zero-shot methods cannot meet precision agriculture requirements

## Practical Contributions

- **Minimum dataset requirements established**: 60-130 samples depending on architecture
- **Predictable performance scaling**: Logarithmic relationship enables resource planning
- **Hybrid approach potential**: Handcrafted methods valuable for